

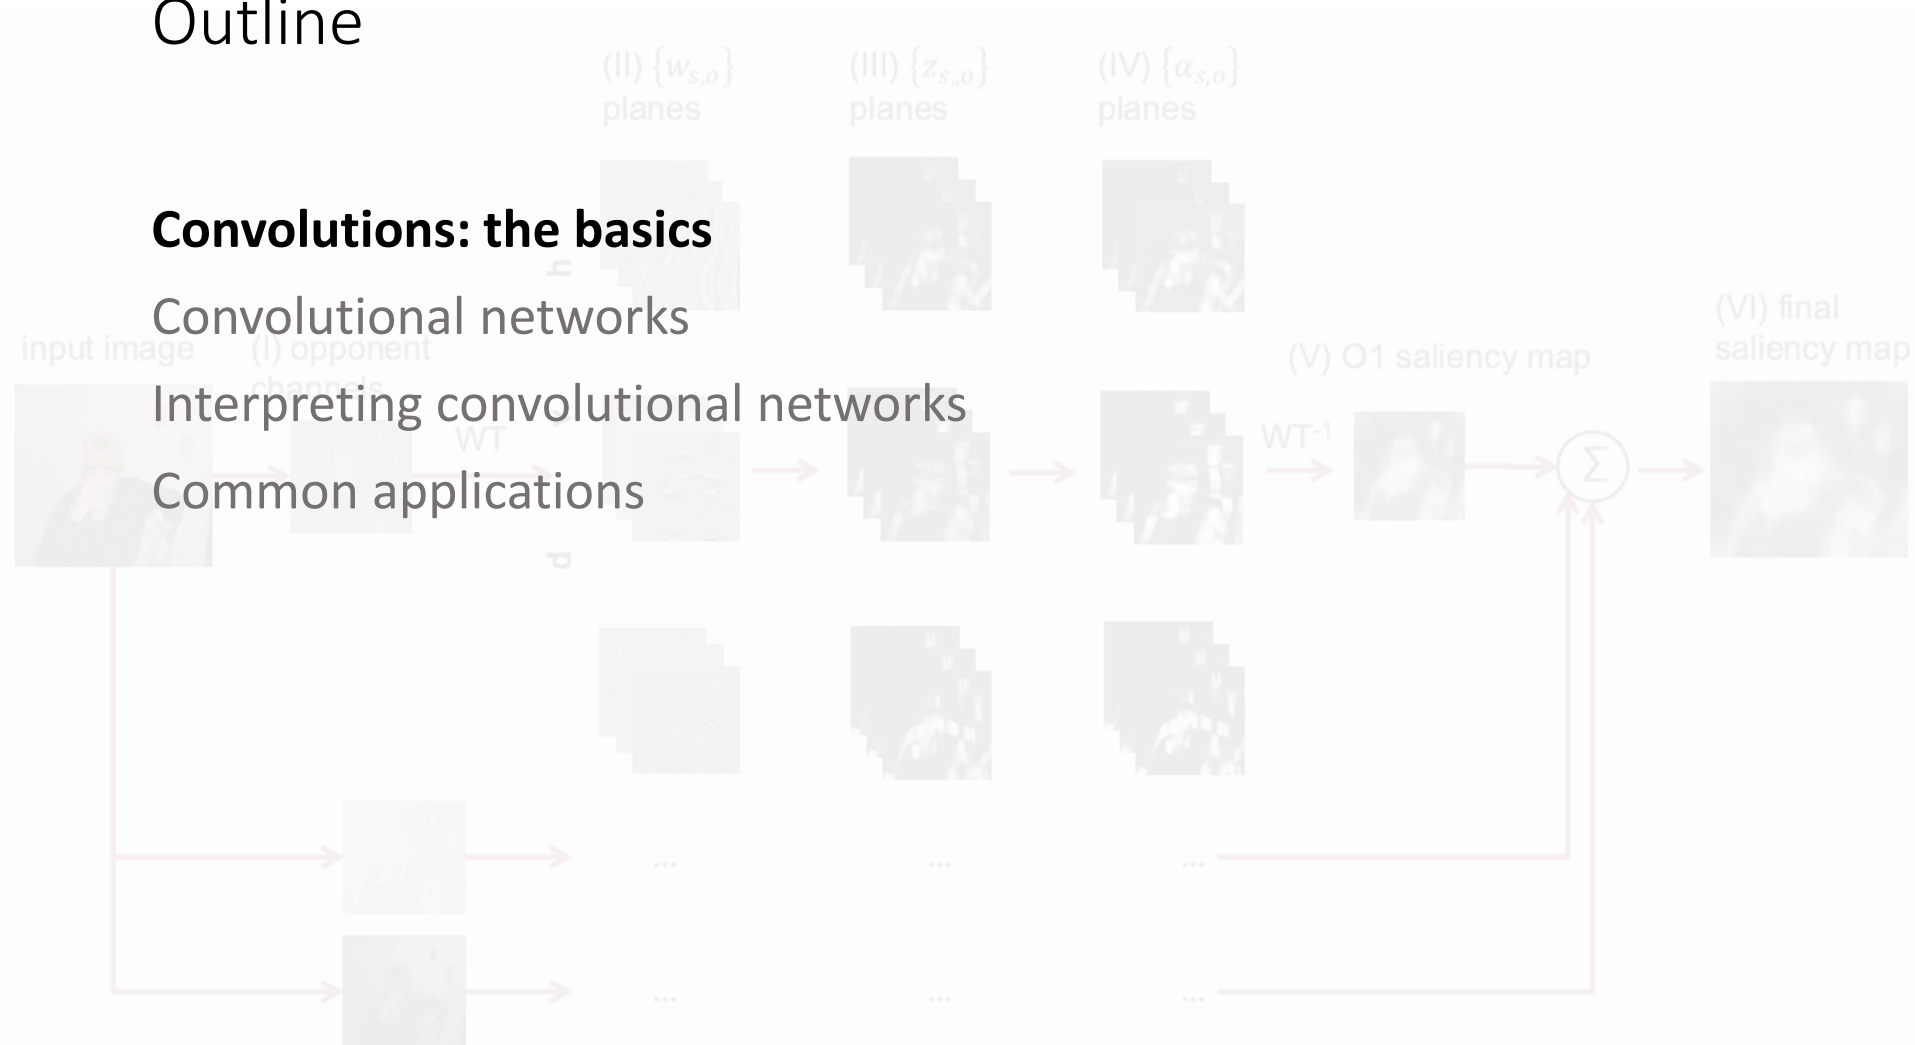
# Outline

## Convolutions: the basics

Convolutional networks

Interpreting convolutional networks

Common applications



# Computer Vision Problems

## Image Classification




64x64

→ Cat? (0/1)

## Neural Style Transfer



## Object detection



# Deep Learning on large images




$64 \times 64 \times 3 = 12.288$  pixels


→ Cat? (0/1)




$1000 \times 1000 \times 3 = 3 \text{ M}$  pixels



# Computer Vision Problem




Detection of layers of Neural Networks



# Convolutional Neural Network

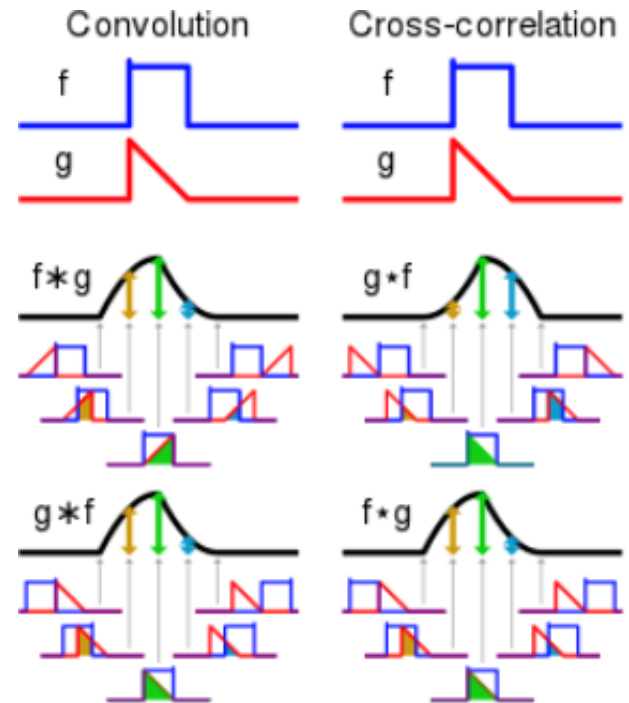
Neural networks that include convolution operations

Training set



# Convolution Layer

- A network layer that convolves its receptive input before passing it to the next layer.
- Most ML libraries implement convolutional layers as **cross-correlation layers**.



$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example

$$(I \star K)(0,0) = I(0,0)K(0,0) + I(0,1)K(0,1) + I(0,2)K(0,2) + \\ I(1,0)K(1,0) + I(1,1)K(1,1) + I(1,2)K(1,2) + \\ I(2,0)K(2,0) + I(2,1)K(2,1) + I(2,2)K(2,2) +$$

$I$

1	0	0	1	2
0	0	0	3	0
0	1	2	1	1
1	1	3	0	0
3	0	0	0	1

$K$

0	0	1
0	2	0
1	1	0


$\star$

$$(K \star I)(0,0) = 1 \cdot 0 + 0 \cdot 0 + 0 \cdot 1 + \\ 0 \cdot 0 + 0 \cdot 2 + 0 \cdot 0 + \\ 0 \cdot 1 + 1 \cdot 1 + 2 \cdot 0 + \\ = 1$$

$$(I \star K)(i,j) == \sum_m \sum_n I(m,n)K(i+m,j+n)$$

# Discrete cross-correlation: 2-D example

Can be viewed as a “sliding window” operation:




$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$



# Discrete cross-correlation: 2-D example


Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example


Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example


Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example


Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example


Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example


Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) = \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example


Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example

Can be viewed as a “sliding window” operation:




$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$



# Discrete cross-correlation: 2-D example

Can be viewed as a “sliding window” operation:



$$(I \star K)(i, j) = \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Discrete cross-correlation: 2-D example

1	0	0	1	2
0	0	0	3	0
0	1	2	1	1
1	1	3	0	0
3	0	0	0	1

★

0	0	1
0	2	0
1	1	0


=

1	4	11
4	11	5
7	7	1


Often called the  
feature map

$$(I \star K)(i, j) = \sum_m \sum_n I(m, n) K(i + m, j + n)$$


# Discrete cross-correlation: 2-D example



★



=



Often called the  
feature map

$$(I \star K)(i, j) = \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Vertical edge detection


10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0




\*




# Vertical edge detection examples

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0




\*

1	0	-1
1	0	-1
1	0	-1




=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0




0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10




\*

1	0	-1
1	0	-1
1	0	-1



=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0



# Vertical and Horizontal Edge Detection

1	0	-1
1	0	-1
1	0	-1

Vertical

1	1	1
0	0	0
-1	-1	-1

Horizontal

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

\*

1	1	1
0	0	0
-1	-1	-1

=

0	0	0	0
30	10	-10	-30
30	10	-10	-30
0	0	0	0

# Discrete cross-correlation: 2-D example

1	0	0	1	2
0	0	0	3	0
0	1	2	1	1
1	1	3	0	0
3	0	0	0	1

★

0	0	1
0	2	0
1	1	0

=

1	4	11
4	11	5
7	7	1

Something's not quite right...

$$(I \star K)(i, j) == \sum_m \sum_n I(m, n) K(i + m, j + n)$$

# Cross-correlation: padding

Adding extra pixels outside the image

0	0	0	0	0	0
0	35	19	25	6	0
0	13	22	16	53	0
0	4	3	7	10	0
0	9	8	1	3	0
0	0	0	0	0	0



# Discrete cross-correlation: padding

0 <sub>0</sub>	0 <sub>0</sub>	0 <sub>1</sub>	0	0	0	0	0	0
0 <sub>0</sub>	0 <sub>2</sub>	0 <sub>0</sub>	0	0	0	0	0	0
0 <sub>1</sub>	0 <sub>1</sub>	1 <sub>0</sub>	0	0	1	2	0	0
0	0	0	0	0	3	0	0	0
0	0	0	1	2	1	1	0	0
0	0	1	1	3	0	0	0	0
0	0	3	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Technically, our signals have infinite extent...  
we solve by padding with zeros

		1	4	11				
		4	11	5				
		7	7	7				

# Discrete cross-correlation: padding

0 <sub>0</sub>	0 <sub>0</sub>	0 <sub>1</sub>	0	0	0	0	0	0
0 <sub>0</sub>	0 <sub>2</sub>	0 <sub>0</sub>	0	0	0	0	0	0
0 <sub>1</sub>	0 <sub>1</sub>	1 <sub>0</sub>	0	0	1	2	0	0
0	0	0	0	0	3	0	0	0
0	0	0	1	2	1	1	0	0
0	0	1	1	3	0	0	0	0
0	0	3	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

0						
		1	4	11		
		4	11	5		
		7	7	1		

# Discrete cross-correlation: padding

0	0 <sub>0</sub>	0 <sub>0</sub>	0 <sub>1</sub>	0	0	0	0	0
0	0 <sub>0</sub>	0 <sub>2</sub>	0 <sub>0</sub>	0	0	0	0	0
0	0 <sub>1</sub>	1 <sub>1</sub>	0 <sub>0</sub>	0	1	2	0	0
0	0	0	0	0	3	0	0	0
0	0	0	1	2	1	1	0	0
0	0	1	1	3	0	0	0	0
0	0	3	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

0	1							
		1	4	11				
		4	11	5				
		7	7	1				

# Discrete cross-correlation: padding

0	0	0 <sub>0</sub>	0 <sub>0</sub>	0 <sub>1</sub>	0	0	0	0
0	0	0 <sub>0</sub>	0 <sub>2</sub>	0 <sub>0</sub>	0	0	0	0
0	0	1 <sub>1</sub>	0 <sub>1</sub>	0 <sub>0</sub>	1	2	0	0
0	0	0	0	0	3	0	0	0
0	0	0	1	2	1	1	0	0
0	0	1	1	3	0	0	0	0
0	0	3	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

0	1	1				
		1	4	11		
		4	11	5		
		7	7	1		

# Discrete cross-correlation: padding

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	1	0	0	1	2	0	0
0	0	0	0	0	3	0	0	0
0	0	0	1	2	1	1	0	0
0	0	1	1	3	0	0	0	0
0	0	3	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Normally we want the output to maintain the input size

0	1	1	0	1	3	2
0	2	0	0	5	7	0
1	0	1	4	11	2	1
0	1	4	11	5	2	0
0	6	7	7	1	1	1
1	7	3	0	0	2	0
3	0	0	0	1	0	0

# Discrete cross-correlation: padding

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	1	0	0	1	2	0	0
0	0	0	0	0	3	0	0	0
0	0	0	1	2	1	1	0	0
0	0	1	1	3	0	0	0	0
0	0	3	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Normally we want the output to maintain the input size

0	1	1	0	1	3	2
0	2	0	0	5	7	0
1	0	1	4	11	2	1
0	1	4	11	5	2	0
0	6	7	7	1	1	1
1	7	3	0	0	2	0
3	0	0	0	1	0	0

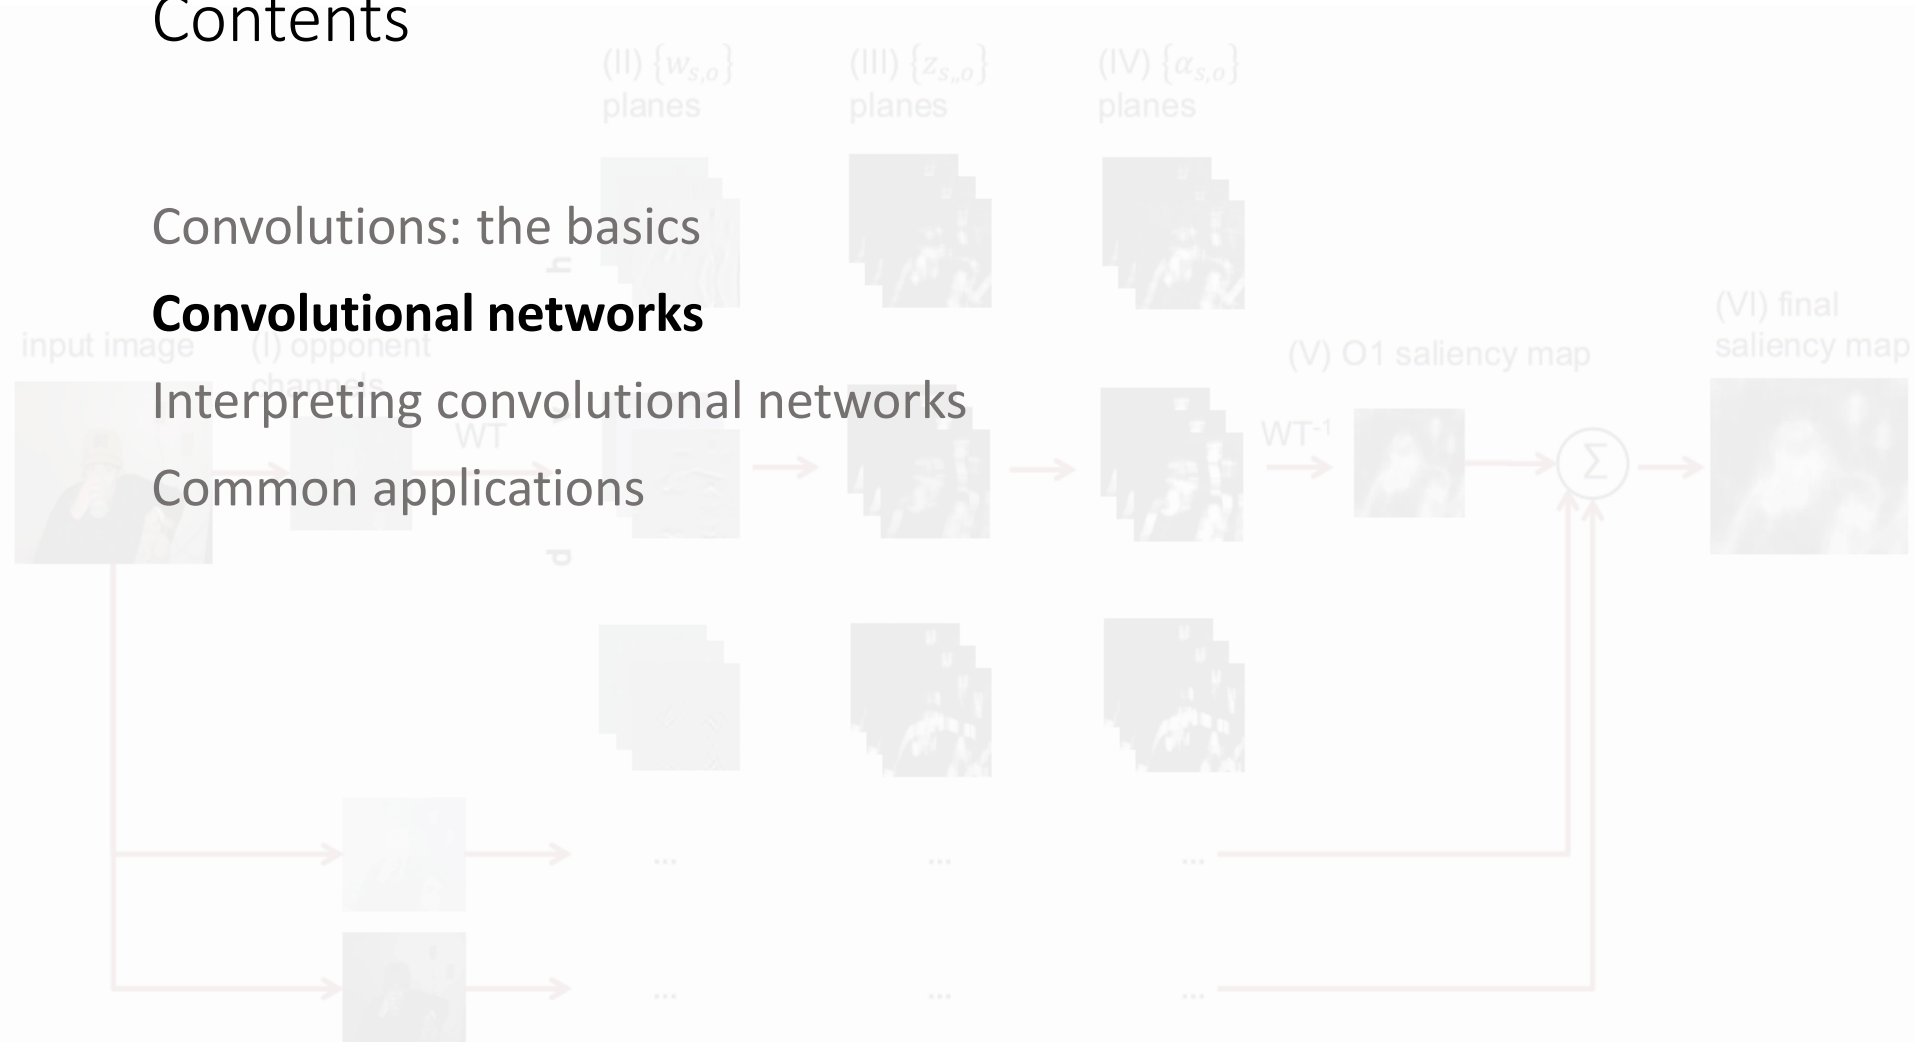
# Contents

Convolutions: the basics

**Convolutional networks**

Interpreting convolutional networks

Common applications



# Convolutional networks

Reminder: Neural networks that include convolution operations

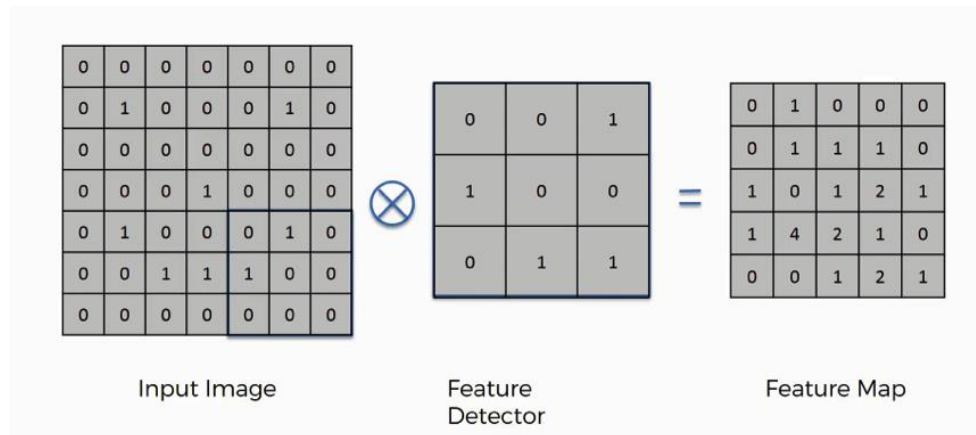
Can be used in place of dense matrix multiplication (i.e. fully-connected layers)

Motivations:

- Sparse connectivity
- Parameter sharing
- Translation equivariance
- Arbitrary input sizes



# Why convolutions: Motivation




**Sparsity of connections:** In each layer, each output value depends only on a small number of inputs.

**Parameter sharing:** A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.

**Translation Equivariance**


**Arbitrary Input Sizes**

# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$

# Neural networks $\Rightarrow$ Convolutional networks




# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$

# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$

# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$

# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$

# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$




$$h_i = \sum_j w_j x_{j+i} + b$$



# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$




$$h_i = \sum_j w_j x_{j+i} + b$$

# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$



$$h_i = \sum_j w_j x_{j+i} + b$$

# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$

dense connectivity vs **sparse connectivity**


In our example:

5x5 multiplications vs **5x3 multiplications**

Sparse connectivity scales better  
e.g. 25x25 vs 25x3



# Neural networks $\Rightarrow$ Convolutional networks




$$\mathbf{h} = \mathbf{W}\mathbf{x} + \mathbf{b}; h_i = \sum_j w_{ij}x_j + b_i$$



unshared vs **shared weights**

In our example:  
5x5 vs **3 weights**

shared weights scale way better:  
e.g. 25x25 vs **3**




# Translation equivariance




# Translation equivariance

Why is translation equivariance useful?


Helps make predictions **translation invariant**



“butterfly”



“butterfly”




“butterfly”


What other types of invariance could be useful?

# Arbitrary input sizes

$$h_i = \sum_j W_{ij} x_j + b_i$$




$$h_i = \sum_j w_j x_{j+i} + b$$




# Arbitrary input sizes

$$h_i = \sum_j W_{ij} x_j + b_i$$



$$h_i = \sum_j w_j x_{j+i} + b$$



- Sparse connectivity
- Parameter sharing
- Translation equivariance
- Arbitrary input sizes



# Strided convolution

2	3	3	4	7	3	4	4	6	3	2	4	9	4
6	1	6	0	9	1	8	0	7	1	4	0	3	2
3	3	4	4	8	3	3	4	8	3	9	4	7	4
7	1	8	0	3	1	6	0	6	1	3	0	4	2
4	3	2	4	1	3	8	4	3	3	4	4	6	4
3	1	2	0	4	1	1	0	9	1	8	0	3	2
0	-1	1	0	3	-1	9	0	2	-1	1	0	4	3

\*

3	4	4
1	0	2
-1	0	3

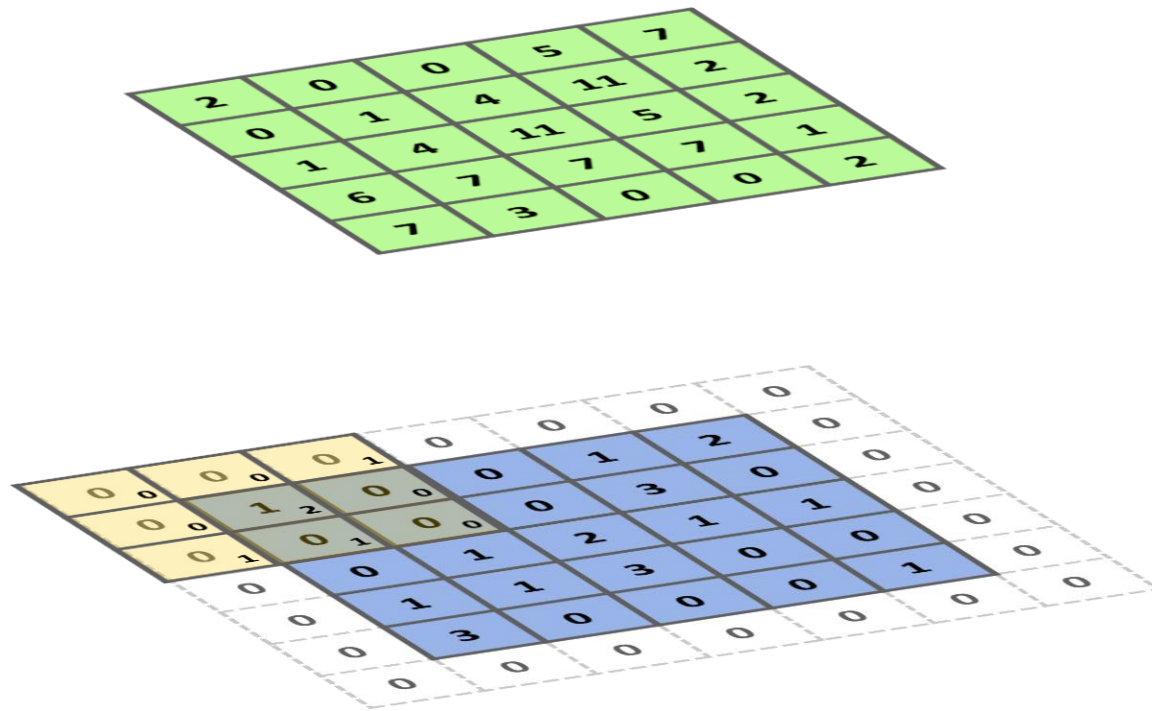
=

91	100	83
69	91	127
44	72	74

Stride = 2


**Stride** is the number of pixels shifts over the input matrix (sliding step).

# The effect of different strides





# The effect of different strides

Stride: 1x1





# The effect of different strides

Stride: 1x1





# The effect of different strides

Stride: 1x1





# The effect of different strides

Stride: 1x1





# The effect of different strides

Stride: 1x1



# The effect of different strides



Stride: 2x2







# The effect of different strides

Stride: 2x2





# The effect of different strides

Stride: 2x2





# The effect of different strides

Stride: 2x2





# The effect of different strides

Stride: 2x2




# The effect of different strides

Stride: 2x2




# The effect of different strides

Stride: 2x2



Why use stride > 1?

- Reduce redundancy
- Compress feature map



# Summary of convolutions


$n \times n$  image       $f \times f$  filter

padding  $p$       stride  $s$

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$

# Pooling


Operation to aggregate or “summarize” sub-region of input.





# Pooling

Operation to aggregate or “summarize” sub-region of input.




$$p = f(\mathbf{a})$$

$$f_{\max}(\mathbf{a}) = \max_i(a_i)$$

$$f_{\text{avg}}(\mathbf{a}) = \frac{1}{N} \sum_{i=1}^N a_i$$

# Pooling

Operation to aggregate or “summarize” sub-region of input.



$$p = f(\mathbf{a})$$

$$f_{\max}(\mathbf{a}) = \max_i(a_i)$$


$$f_{\text{avg}}(\mathbf{a}) = \frac{1}{N} \sum_{i=1}^N a_i$$

# Pooling

Operation to aggregate or “summarize” sub-region of input.

Shift input by 1 position:

5/5 inputs change but only 1/3 pooled outputs



$$p = f(\mathbf{a})$$

$$f_{\max}(\mathbf{a}) = \max_i(a_i)$$


$$f_{\text{avg}}(\mathbf{a}) = \frac{1}{N} \sum_{i=1}^N a_i$$

# Pooling

Operation to aggregate or “summarize” sub-region of input.

Shift input by 1 position:

5/5 inputs change but only 1/3 pooled outputs



$$p = f(\mathbf{a})$$

$$f_{\max}(\mathbf{a}) = \max_i(a_i)$$

$$f_{\text{avg}}(\mathbf{a}) = \frac{1}{N} \sum_{i=1}^N a_i$$

# Pooling


Operation to aggregate or “summarize” sub-region of input.

Shift input by 1 position:

5/5 inputs change but only 1/3 pooled outputs

Why is pooling useful?

- Adds **translation invariance** (useful when we care about “what” more than “where”)
- Can be used to **compress signal** (useful to improve computational efficiency)



$$p = f(\mathbf{a})$$

$$f_{\max}(\mathbf{a}) = \max_i(a_i)$$

$$f_{\text{avg}}(\mathbf{a}) = \frac{1}{N} \sum_{i=1}^N a_i$$

# Pooling


Operation to aggregate or “summarize” sub-region of input.

Shift input by 1 position:

5/5 inputs change but only 1/3 pooled outputs

Why is pooling useful?

- Adds **translation invariance** (useful when we care about “what” more than “where”)
- Can be used to **compress signal** (useful to improve computational efficiency)



$$p = f(\mathbf{a})$$

$$f_{\max}(\mathbf{a}) = \max_i(a_i)$$

$$f_{\text{avg}}(\mathbf{a}) = \frac{1}{N} \sum_{i=1}^N a_i$$


## Receptive fields

Borrowed terminology from neuroscience:

⇒ stimulus region that impacts a neuron's firing

For neural networks:

⇒ Region of input signal that impacts node output



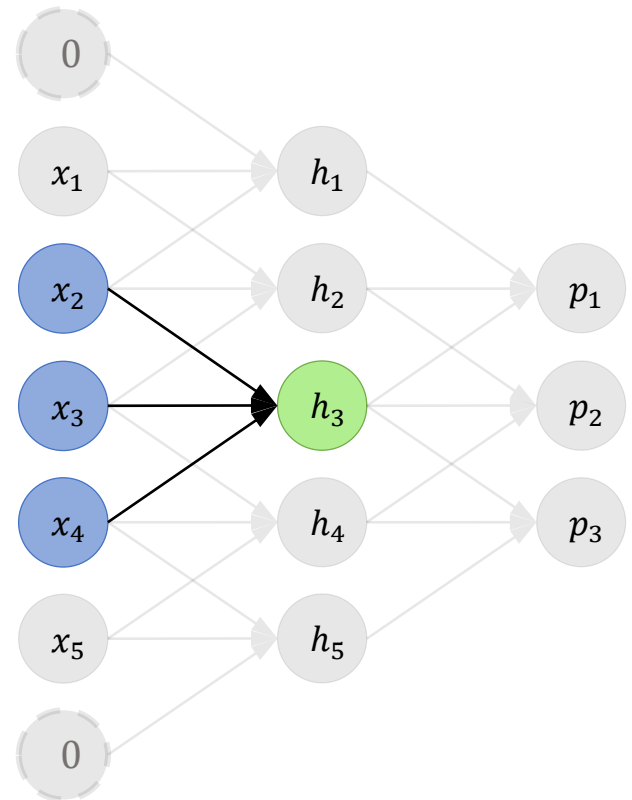
## Receptive fields

Borrowed terminology from neuroscience:

⇒ stimulus region that impacts neuronal firing

For neural networks:

⇒ Region of input signal that impacts node output





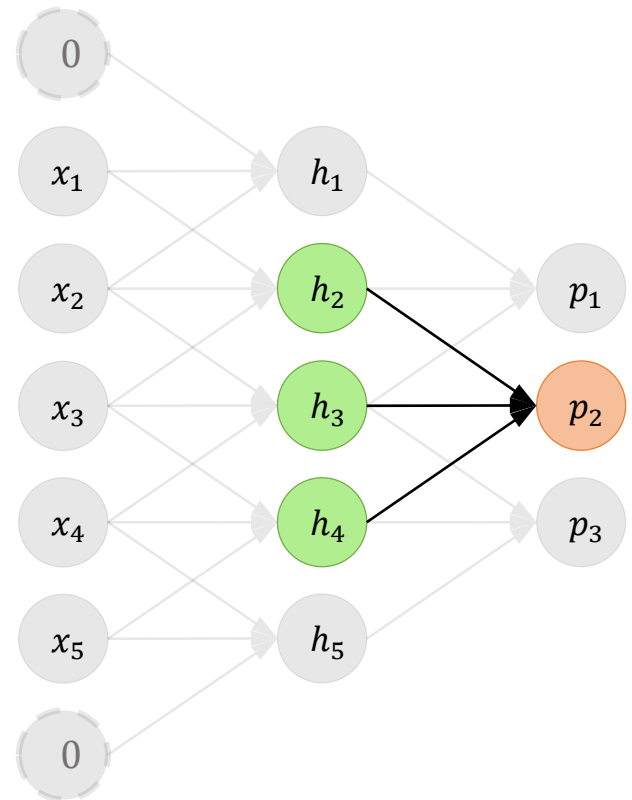
## Receptive fields

Borrowed terminology from neuroscience:

⇒ stimulus region that impacts neuronal firing

For neural networks:

⇒ Region of input signal that impacts node output



## Receptive fields

Borrowed terminology from neuroscience:


⇒ stimulus region that impacts neuronal firing

For neural networks:


⇒ Region of input signal that impacts node output

**Effective** field size:


⇒ Region of **network** input signal that impacts node output



# Convolutions on RGB image




# Multiple filters



6 x 6 x 3


\*

Vertical Edge




3 x 3 x 3

=




4 x 4

\*



3 x 3 x 3

=



4 x 4

Horizontal Edge